CheckoutGuard — Business plan

Prepared 11 October 2026. USD. Pre-revenue, assumption-led plan; not validated traction.

Decision summary

Recommend a capped Team paid pilot, not unlimited usage or automated purchasing. Known catalog-priced token estimate: $0.00315467/successful synthetic run; mean supplied runtime: 22.2933 seconds. Modeled delivery cost including provisional retry, browser, database, evidence, telemetry and run support allowances: $0.0190856/completed run. Neither is a confirmed charge.

Base break-even: approximately 96 subscriptions at the assumed plan mix, or $8,650.67 monthly revenue. Fixed costs of $6,800/month include a $6,000 founder-pay assumption. Base first nonnegative operating month: March 2028; peak accumulated operating loss before that is approximately $59,175. This is a model-implied operating deficit, not verified funding need, cash balance or runway. Collections, tax, capex and liabilities remain unknown.

Proposed pricing — assumptions

Plan	Monthly price	Included workspaces	Pooled completed runs	Contribution at 50% use	Margin at full quota
Developer	$29	1	100	$21.88	72.1%
Team	$99	3	500	$78.46	74.4%
Agency	$299	10	2,000	$240.64	74.1%

Support and fees are assumptions; the margins do not include fixed overhead.

Revenue scenarios — all assumptions

Case	Month 12 revenue (Oct 2027)	Month 24 revenue (Oct 2028)	Month 24 subscriptions (expected)	Month 24 operating result	First nonnegative operating month
Conservative	$1,864	$2,871	31.9	$-4,527	Not reached through October 2028
Base	$5,586	$13,907	154.3	$4,209	2028-03
Growth	$12,342	$40,765	458.0	$25,464	2027-08

No implied actual users or sales. Fractional subscriptions represent expected counts. Month 24 revenue is monthly, not annual revenue.

Known supplied runtime observations

Synthetic run	Decisions	Screenshots	Seconds	Input tokens	Output tokens	Catalog estimate
1	7	8	21.30	12,227	202	$0.00315775
2	7	8	28.92	12,128	179	$0.00312150
3	7	8	16.66	12,335	202	$0.00318475

Evidence.md was referenced but not attached. No independent runtime verification is claimed. Configured request: Apertus; response label: Qwen SEA-LION v4 32B IT. Identity not independently verified. Earlier timed-out request usage remains unknown.

Cost formula

Cost/completed run = known token estimate × provider price multiplier × (1 + extra-attempt ratio × extra-attempt token allowance) + browser seconds/attempt ÷ 3,600 × browser/hour × (1 + extra-attempt ratio) + incremental database/run + evidence/run + telemetry/run + support minutes/run ÷ 60 × support/hour.

Contribution = subscription revenue − completed runs × cost/run − workspace support − assumed payment fees. Operating result = contribution − fixed monthly costs.

One-way sensitivity — base October 2028

Test	Monthly operating result	Change from base
Operating result	$4,209	$0
Provider 10x operating result	$2,923	$-1,285
Extra attempts 100 percent operating result	$4,068	$-140
Browser 3x operating result	$4,073	$-136
Support 2x operating result	$2,117	$-2,092
Full quota usage operating result	$3,489	$-720

Provider pricing ×10 reduces that month’s operating result by about $1,285; doubling all support reduces it by about $2,092. Do not equate cheap token usage with a cheap service.

CheckoutGuard — Business plan

CheckoutGuard

Pre-revenue developer checkout QA — assumption-led business plan, 11 October 2026.

Goal-driven checkout QA for developers and ecommerce agencies. Prototype: Go, React/TypeScript, Redux Toolkit, PostgreSQL and Chrome. Checks totals against integer price breakdowns, confirms stored synthetic orders, keeps screenshots and compares corrected retests.

Evidence boundary: no validated paying customers or revenue. All pricing, customer counts, conversion, churn and revenue projections are assumptions—not traction. Runtime observations below are supplied by the user; EVIDENCE.md has not been attached. No funding, cash balance, runway, provider invoice or production purchase authority is established.

Safety boundary: owned or explicitly authorized staging environments and synthetic orders only. Experimental Shopee/Tokopedia adapters prepare checkout and stop before final order confirmation. No automatic goods purchases or operation of merchant payment systems is assumed.

Strategy: checkout-specific QA, not a general purchasing agent

Buyer: agency technical lead, ecommerce developer or QA owner responsible for checkout regressions. Start with owned or authorized staging sites and seeded synthetic orders. Core value: integer-safe totals assertions, persisted-order checks, screenshot evidence and corrected-retest comparisons. Alternative workflows are manual checklists, Playwright/Cypress scripts and general browser agents; CheckoutGuard must prove less setup and clearer fault evidence than these alternatives, not just that it can click through a checkout.

MVP priorities: reproducible fixtures, environment allowlists, synthetic test identities, idempotency and cleanup, correlation IDs, CI artifact links, timeout/retry budgets, model-routing/usage visibility and deletion controls. React/TypeScript and Redux manage workflow; Go orchestrates browser checks and PostgreSQL stores run/order assertions. Do not promise production merchant access. Experimental marketplace adapters are outside the initial paid offering.

Distribution: recruit design partners through ecommerce developer communities and agency leads; run founder-led demos on their staging checkout, then integrate into CI. Budgeted sales spend is an unvalidated capacity assumption. No validated demand, sales, funding or users are claimed.

Validation plan — proposed gates, not results

Weeks 1–2: interview 15 qualified developers/agency QA owners about their last checkout incident, current scripts, time cost, approval/security rules and evidence needs. Gate: at least 8 describe a recent recurring problem and 5 grant access to an authorized staging test. Record refusals as well as enthusiasm.

Weeks 3–4: 5 design partners run at least 20 representative cases each, including discounts, taxes, shipping, currency rounding, broken totals, persisted-order failures and corrected retests. Instrument completion, false positives/negatives, tokens per attempt, timeout billing, screenshots/bytes, browser time and support minutes. Proposed gates: ≥95% task completion on supported fixtures, ≥90% seeded-defect detection, ≤5% false alarms, and p95 completion ≤60 seconds. These targets are not established by three synthetic runs. Human review confirms all seeded outcomes.

Weeks 5–6: present real paid offers—not only willingness-to-pay surveys. Test Team $99 with 500 pooled runs against $149 with the same quota and an Agency $299 offer. Aim for 3 paid pilots from 10 qualified trial organizations, usage in at least 3 of 4 weeks and two renewal commitments after a full month. No sale is recorded until payment is real. Stop expanding acquisition if activation/renewal fail; diagnose trust, integration effort, false alarms and security objections first.

Weeks 7–8: reconcile invoices and run logs; test p95 high-token/retry cases, support and full-quota margins. Gate: ≥70% contribution margin at the intended mix, sustainable support burden, and documented authorization/data-retention controls. Replace assumptions in this model with observed cohort figures before claiming a repeatable business.

Runway — deliberately uncalculated

Simple steady-burn runway = unrestricted available cash ÷ monthly net cash burn, only when burn is positive. A better answer is the first month when opening unrestricted cash + cumulative actual cash receipts − cash payments + committed financing falls below minimum cash. This plan’s cumulative operating result is not a cash balance or a runway claim.

Needed: cash amount and as-of date/currency; restricted cash; minimum cash buffer; founder pay timing; actual hosting/provider/payment invoices; taxes; collections/refunds; payables; capex; debt service; and only documented committed financing. Variable burn requires a dated cash-flow schedule, not dividing by one optimistic month. No funding, balance or runway has been invented.

CheckoutGuard — Assumptions

Proposed prices — test, don’t promise

All values on this page are editable planning assumptions in USD, excluding tax. Developer: $29/month, one workspace, 100 completed runs. Team: $99/month, three workspaces, 500 runs pooled. Agency: $299/month, ten workspaces, 2,000 runs pooled. Charge per organization

subscription, not per developer seat. Start with Team-focused paid pilots; do not launch unlimited usage. No overage revenue is modeled: default to a hard cap plus explicit upgrade approval; if validated, test $0.10/completed-run overages with a monthly spending cap. Corrected retests consume a run. Provider retries should not consume a second customer quota unit.

Prices buy auditable checkout checks, regression evidence and less QA work—not tokens. Retention assumption: 30 days for screenshots, longer only by explicit agreement. The $0.001 evidence/run allowance is provisional: 8 screenshots × assumed 0.25 MB × retention and storage/request rates plus export/egress; measure real files before contracting.

Cost basis and unknowns

Public AI catalog lists Qwen SEA-LION v4 32B IT at $0.25/input million and $0.50/output million: https://platform.publicai.co/models . Checked 11 October 2026. Requested Apertus and response labels disagree; model identity is not independently verified. These rates estimate known reported usage; they do not establish actual model routing or charges.

Three successful synthetic runs are not a workload distribution or reliability estimate. Earlier timeout token usage is unknown. **Extra attempt ratio = additional attempts per completed run**, assumed 20%, not measured failure probability. Timeout token allowance ratio assumes each extra attempt uses 100% of a successful run’s tokens; test 0.5×–2× and obtain provider logs/invoices. Never replace unknown timeout usage with zero. Browser assumption is 45 seconds/attempt at $0.12/hour, above the supplied 22.29-second mean to allow cold-start/setup time. Browser/Go worker cost is included in that rate; fixed React/API hosting is separate.

DB, storage, telemetry and support costs are placeholders, not invoices. Database $0.002/completed run and telemetry $0.0005/run are incremental allowances above the separate fixed service baselines. Support has both 0.02 minutes/run and 10/25/60 minutes/subscription/month at $30/hour, as incremental capacity cost. Founder compensation $6,000/month covers development/product time, excluding that support capacity; avoid counting the same hours twice. Other fixed costs: hosting $150, DB $60, telemetry $30, overhead $60, sales $500/month: total $6,800 including founder compensation.

Acquisition and forecast conventions

Base qualified prospects/month at launch: Developer 40, Team 20, Agency 5; assumed conversion 5%/10%/10%; pipeline grows 5% monthly and churn is 3% monthly. Conservative: conversion 2.5%/5%/5%, no pipeline growth, 5% churn. Growth: conversion 10%/15%/20%, 8% monthly pipeline growth, 2% churn. These are explicit targets, not observed lead flow, conversion or retention.

October 2026 is a zero-subscription **planning seed**, not a historical customer census; launch is assumed November 2026, with 24 forecast months through October 2028. Fractions are expected subscriptions, not actual partial customers. Revenue uses full-month ending subscriptions, collected in that month by assumption; no annual prepay, expansion, refunds,

taxes or receivable delay are included. Payment processing uses an unverified 3% + $0.30/subscription/month assumption. Customer mix changes only through acquisition and churn. No financing or cash balance is modeled.

CheckoutGuard — Forecast

Reading this forecast

Main Scenario is the base case. Conservative and Growth change only assumed acquisition conversion, monthly prospect growth and churn; pricing, usage, support and fixed costs stay the same. New subscriptions = qualified prospects × conversion; lost subscriptions = prior ending subscriptions × churn; ending = prior + new − lost. Ending subscriptions receive a full month of revenue, an optimistic billing convention that should be replaced with real subscription start dates.

Contribution subtracts browser/model delivery, incremental database/evidence/telemetry, run-level and workspace support, and **assumed payment fees of 3% + $0.30/subscription/month**. They are not verified merchant/provider charges. Fixed $6,800/month is then deducted once to give operating result, before tax, depreciation and financing.

Break-even subscriptions = ceiling(fixed monthly cost ÷ weighted contribution/subscription). Break-even revenue uses that same customer mix. It is a steady-state threshold, not proof that customers can be acquired. Cumulative operating result starts from zero as a profit accumulation seed, not a cash balance. Cash burn may differ with collections, prepayments, tax, capex and liabilities.

One-way sensitivities hold customer demand and prices fixed: model token prices ×10; extra attempts increase from 20% to 100% per completed run; browser compute ×3; all support ×2; or run quota usage 50%→100%. Effects show in operating result, not confirmed spending. Provider cost and support changes do not silently change conversion or churn. A usage increase requires delivery capacity and may affect service quality; those second-order effects are not modeled.

Deliverable verification and limits

- base: all 25 periods matched independent calculations for subscriptions, revenue, contribution, operating result and cumulative result.

- conservative: all 25 periods matched independent calculations for subscriptions, revenue, contribution, operating result and cumulative result.

- growth: all 25 periods matched independent calculations for subscriptions, revenue, contribution, operating result and cumulative result.

- Five one-way sensitivities matched independent October 2028 calculations; signs reduce operating results as expected.

- Annual native reads checked: flows sum monthly values; 2028 year-end balance is intentionally absent because the horizon stops in October. No full-year 2028 balance is asserted.

- Actual export: scanned XML/relationships for personal/account/credential identifiers, checked external links and Excel error cells. No matched private identifiers or external workbook links found; export preserved byte-for-byte.

- Native model simulator rendered without reported errors or warnings in light and dark at 1,200px width. Visible screenshots were reviewed in chat. Controls/click actions and narrower widths were not browser-tested.

- The reviewed Canvas screenshots were not returned as downloadable image files by cfo.ai, so they are **not included in this artifact**. Do not claim this package preserves those images. Attach the screenshots from the app/chat to complete that evidence requirement.

- This package is prepared for public review, but has not been published or given a public share link. No account, payment-method, subscription or purchasing operation was performed.